

# Big Data Part One: An Intro to MapReduce

Chris Jermaine

Rice University

# 15 Years Ago...

- Say you had a big data set, wanted platform to analyze it
- What is “big”?
  - ▶ Too large to fit in RAM of an expensive server machine

# You Might Roll Your Own Software

- Costly, time consuming
  - ▷ A \$10M software feature might eat up most of the IT budget for a single firm
- Requires expertise not always found in house
- Risky: high potential for failure

# Or, You Might Buy a DB System

- Costs a LOT of money
- Performance often unpredictable, or just flat out poor
- Software insanely complicated to use correctly
- Software stack too big/deep, not possible to unbundle
  - ▷ If you are doing analysis, ACID not important
  - ▷ And yet, you pay for it (money, complexity, performance)
- Difficult to put un- or semi-structured data into an SQL DB

# Plus, Many People Just Don't Like SQL

- People uncomfortable with declarative programming
  - ▷ We love it!
  - ▷ But user doesn't really know what's happening under the hood
  - ▷ Makes many programmers uncomfortable
- Also, not easy/natural to specify important computations
  - ▷ Especially data mining and machine learning
  - ▷ Not to mention HPC-style computations

# By Early-Mid 2000's...

- The Internet companies (Google, Yahoo, etc.)...
  - ▷ ...had some of the largest databases in the world
  - ▷ But they never used classical SQL databases for webscale
- How'd they do it?
  - ▷ Many ways...
  - ▷ But paradigm with most widespread impact was MapReduce
  - ▷ First described in a 2004 academic paper, appeared in OSDI
  - ▷ Easy read! Do a search on “Google MapReduce paper”

# What Is MapReduce?

- It is a simple data processing paradigm
- To process a data set:
  - ▷ You have two pieces of user-supplied code
  - ▷ A Map code
  - ▷ And a Reduce code
- These are run in a huge shared-nothing compute cluster
  - ▷ Using three data processing phases
  - ▷ A Map phase
  - ▷ A Shuffle phase moves data around
  - ▷ And a Reduce phase

# First: What Is Shared-Nothing?

- Store/analyze data on a large number of commodity machines
  - ▷ Local, non-shared storage attached to each of them
  - ▷ Only link is via a LAN
  - ▷ Shared nothing refers to no sharing of RAM, storage
  - ▷ Note: NAS is common now, “pure” shared-nothing rarer
- Why good?
  - ▷ Inexpensive, built out of commodity components
  - ▷ Compute resources scales nearly linearly with money
  - ▷ Contrast to shared RAM machine with uniform memory access



# MapReduce: The Map Phase

- Input data are stored in a huge file
  - ▷ Contains a simple list of pairs of type  $(key1, value1)$
- Have a UDF of the form  $Map(key1, value1)$ 
  - ▷ outputs a list of pairs of the form  $(key2, value2)$
- In the Map phase of the MapReduce computation
  - ▷ The *Map* function is called for every record in the input
  - ▷ Instances of *Map* run in parallel all over the cluster

# Example: WordCount

- Large text corpus
- Want to count number of occurs of each word
- Ex output: ('The', 1832321), ('An', 1732432), etc.
- To power the Map phase:
  - ▷ MapReduce software automatically breaks corpus into large number of (*lineNo*, *text*) pairs

# Example Map Phase...


Node 1




Node 3



Node 2



Node 4



Hi mom how are you  
I am fine  
How about you  
We are all good  
Are you good  
I am fine as well  
Not doing too well  
Be better tomorrow

Node 1

(1, Hi mom how are you)  
(5, Are you good mom)

Node 3

(3, How about you dad)  
(7, Not doing too well)

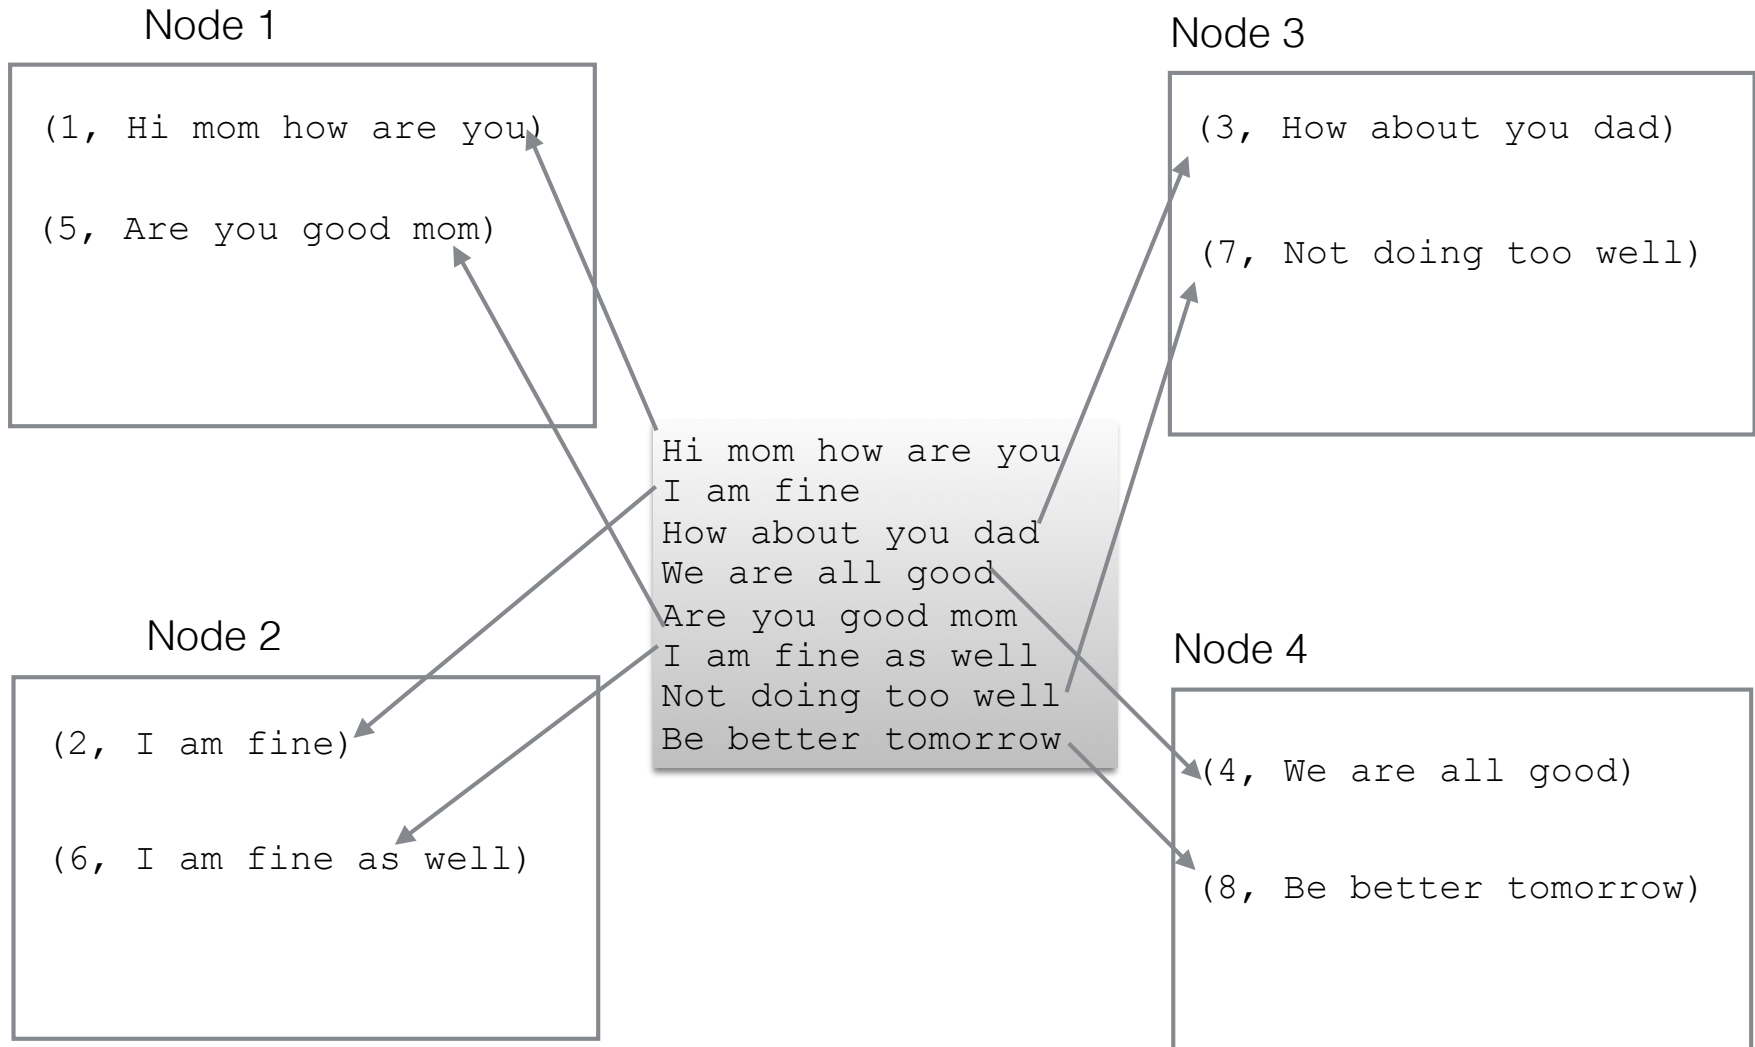
Node 2

(2, I am fine)  
(6, I am fine as well)

Node 4

(4, We are all good)  
(8, Be better tomorrow)

Hi mom how are you  
I am fine  
How about you dad  
We are all good  
Are you good mom  
I am fine as well  
Not doing too well  
Be better tomorrow



### Node 1

(1, Hi mom how are you)  
Apply mapper

(5, Are you good mom)

### Node 2

(2, I am fine)

(6, I am fine as well)

### Node 3

(3, How about you dad)

(7, Not doing too well)

Hi mom how are you  
I am fine  
How about you dad  
We are all good  
Are you good mom  
I am fine as well  
Not doing too well  
Be better tomorrow

### Node 4

(4, We are all good)

(8, Be better tomorrow)

### Node 1

(hi, 1) (mom, 1) (how, 1)  
(are, 1) (you, 1)  
(5, Are you good mom)

### Node 2

(2, I am fine)  
  
(6, I am fine as well)

### Node 3

(3, How about you dad)  
  
(7, Not doing too well)

Hi mom how are you  
I am fine  
How about you dad  
We are all good  
Are you good mom  
I am fine as well  
Not doing too well  
Be better tomorrow

### Node 4

(4, We are all good)  
  
(8, Be better tomorrow)

### Node 1

(hi, 1) (mom, 1) (how, 1)  
(are, 1) (you, 1)

**(5, Are you good mom)**

**Apply mapper**

### Node 2

(2, I am fine)

(6, I am fine as well)

### Node 3

(3, How about you dad)

(7, Not doing too well)

Hi mom how are you  
I am fine  
How about you dad  
We are all good  
Are you good mom  
I am fine as well  
Not doing too well  
Be better tomorrow

### Node 4

(4, We are all good)

(8, Be better tomorrow)

### Node 1

(hi, 1) (mom, 1) (how, 1)  
(are, 1) (you, 1)  
(are, 1) (you, 1)  
(good, 1) (mom, 1)

### Node 2

(2, I am fine)  
  
(6, I am fine as well)

### Node 3

(3, How about you dad)  
  
(7, Not doing too well)

Hi mom how are you  
I am fine  
How about you dad  
We are all good  
Are you good mom  
I am fine as well  
Not doing too well  
Be better tomorrow

### Node 4

(4, We are all good)  
  
(8, Be better tomorrow)



# Done in parallel all over the cluster...

Node 1

```
(hi, 1) (mom, 1) (how, 1)
(are, 1) (you, 1)
(are, 1) (you, 1)
(good, 1) (mom, 1)
```

Node 3

```
(how, 1) (about, 1)
(you, 1) (dad, 1)
(not, 1) (doing, 1)
(too, 1) (well, 1)
```

Node 2

```
(i, 1) (am, 1) (fine, 1)
(i, 1) (am, 1) (fine, 1)
(as, 1) (well, 1)
```

```
Hi mom how are you
I am fine
How about you dad
We are all good
Are you good mom
I am fine as well
Not doing too well
Be better tomorrow
```

Node 4

```
(we, 1) (are, 1)
(all, 1) (good, 1)
(be, 1) (better, 1)
(tomorrow, 1)
```

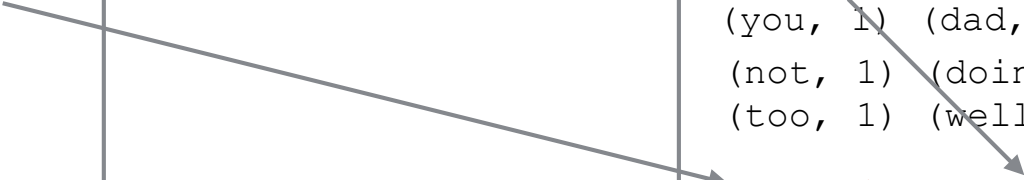
# MapReduce: The Shuffle Phase

- Accepts all of the  $(key2, value2)$  pairs from the Map phase
  - ▷ And it groups them together
- After grouping, all of the pairs
  - ▷ From all over the cluster having the same  $key2$  value
  - ▷ Are merged into a single  $(key2, list\langle value2 \rangle)$  pair
- Called a “Shuffle”...
  - ▷ Because this is where a potential all-to-all data transfer happens

# Example Shuffle Phase...

Node 1

```
(hi, 1) (mom, 1) (how, 1)
(are, 1) (you, 1)
(are, 1) (you, 1)
(good, 1) (mom, 1)
```



Node 3

```
(how, 1) (about, 1)
(you, 1) (dad, 1)
(not, 1) (doing, 1)
(too, 1) (well, 1)
(how, 1) (how, 1)
```

Node 2

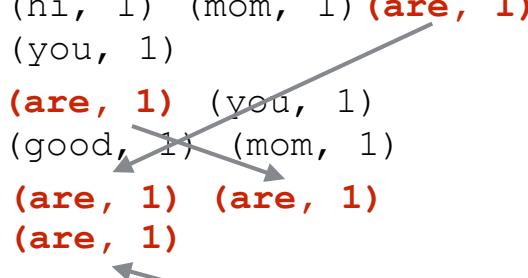
```
(i, 1) (am, 1) (fine, 1)
(i, 1) (am, 1) (fine, 1)
(as, 1) (well, 1)
```

Node 4

```
(we, 1) (are, 1)
(all, 1) (good, 1)
(be, 1) (better, 1)
(tomorrow, 1)
```

Node 1

(hi, 1) (mom, 1) **(are, 1)**  
(you, 1)  
**(are, 1)** (you, 1)  
(good, 1) (mom, 1)  
**(are, 1) (are, 1)**  
**(are, 1)**



Node 3


(about, 1) (you, 1)  
(dad, 1)  
(not, 1) (doing, 1)  
(too, 1) (well, 1)  
**(how, 1) (how, 1)**

Node 2

(i, 1) (am, 1) (fine, 1)  
(i, 1) (am, 1) (fine, 1)  
(as, 1) (well, 1)

Node 4

(we, 1) **(are, 1)**  
(all, 1) (good, 1)  
(be, 1) (better, 1)  
(tomorrow, 1)



Node 1

(hi, 1) (mom, 1) (you, 1)  
(you, 1) (good, 1)  
(mom, 1)  
(are, 1) (are, 1)  
(are, 1)

Node 2

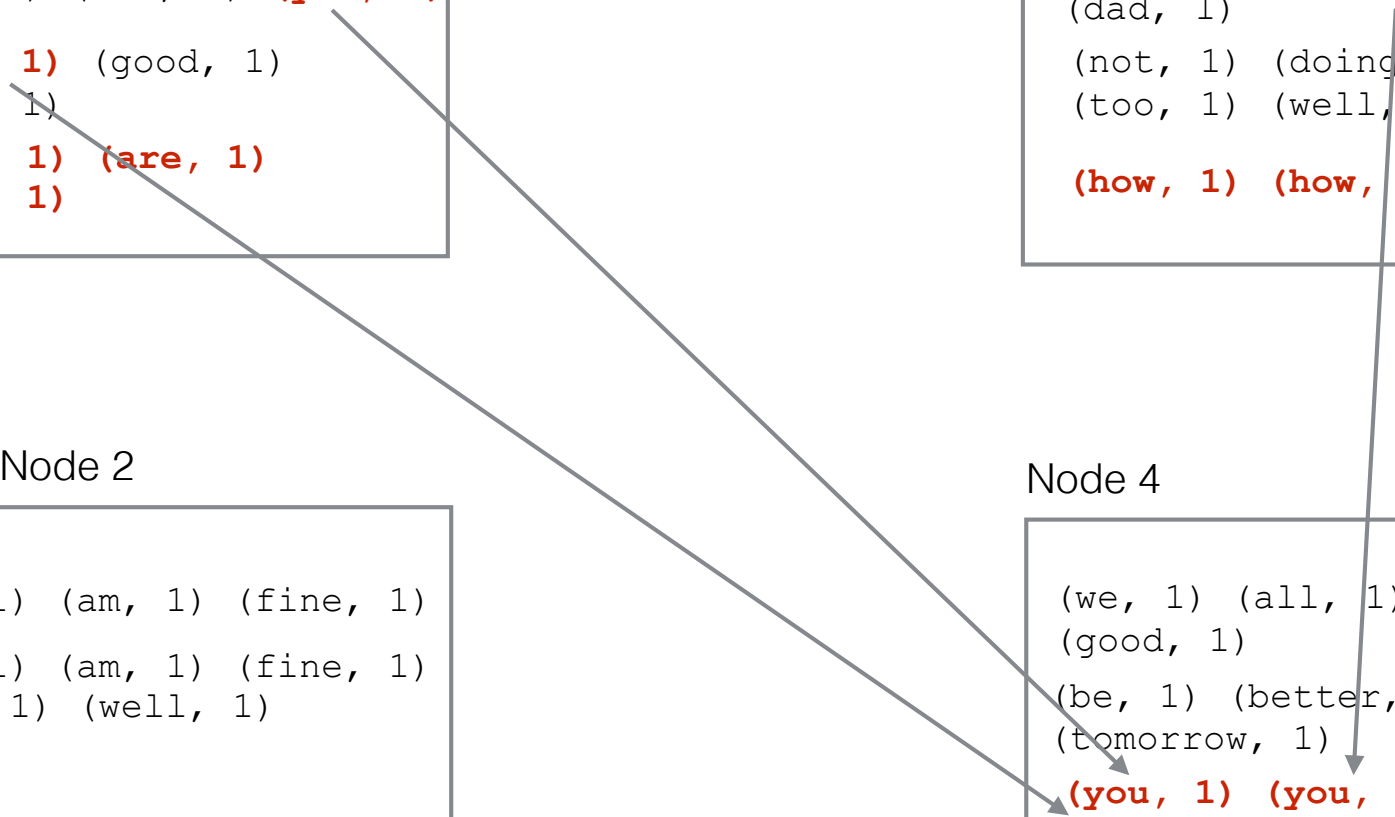
(i, 1) (am, 1) (fine, 1)  
(i, 1) (am, 1) (fine, 1)  
(as, 1) (well, 1)

Node 3

(about, 1) (you, 1)  
(dad, 1)  
(not, 1) (doing, 1)  
(too, 1) (well, 1)  
(how, 1) (how, 1)

Node 4

(we, 1) (all, 1)  
(good, 1)  
(be, 1) (better, 1)  
(tomorrow, 1)  
(you, 1) (you, 1)  
(you, 1)



# After all data have been xferred...

Node 1

```
(am, 1) (am, 1)
(as, 1)
(not, 1)
(good, 1) (good, 1)
(are, 1) (are, 1)
(are, 1)
```

Node 3

```
(better, 1)
(be, 1)
(about, 1)
(too, 1)
(all, 1)
(how, 1) (how, 1)
(hi, 1)
```

Node 2

```
(fine, 1) (fine, 1)
(doing, 1)
(well, 1) (well, 1)
(dad, 1)
(i, 1) (i, 1)
```

Node 4

```
(mom, 1) (mom, 1)
(we, 1)
(tomorrow, 1)
(you, 1) (you, 1)
(you, 1)
```

# Arranged into (key, list) pairs...

Node 1

```
(am, <1, 1>)  
(as, <1>)  
(not, <1>)  
(good, <1, 1>)  
(are, <1, 1, 1>)
```

Node 3

```
(better, <1>)  
(be, <1>)  
(about, <1>)  
(too, <1>)  
(all, <1>)  
(how, <1, 1>)  
(hi, <1>)
```

Node 2

```
(fine, <1, 1>)  
(doing, <1>)  
(well, <1, 1>)  
(dad, <1>)  
(i, <1, 1>)
```

Node 4

```
(mom, <1, 1>)  
(we, <1>)  
(tomorrow, <1>)  
(you, <1, 1, 1>)
```

# MapReduce: The Reduce Phase

- Have a user-supplied function of the form
  - ▷ *Reduce*(*key2*, *list*(*value2*))
  - ▷ Outputs a list of *value3* objects
- In the Reduce phase of the MapReduce computation
  - ▷ *Reduce* function is called for every *key2* value output by the Shuffle
  - ▷ Instances of *Reduce* run in parallel all over the compute cluster
  - ▷ The output of all of those instances is collected
  - ▷ Put in a (potentially) huge output file



# Finally, Reduce Phase...

Node 1

```
(am, <1, 1>)  
(as, <1>)  
(not, <1>)  
(good, <1, 1>)  
(are, <1, 1, 1>)
```

Node 3

```
(better, <1>)  
(be, <1>)  
(about, <1>)  
(too, <1>)  
(all, <1>)  
(how, <1, 1>)  
(hi, <1>)
```

Node 2

```
(fine, <1, 1>)  
(doing, <1>)  
(well, <1, 1>)  
(dad, <1>)  
(i, <1, 1>)
```

Node 4

```
(mom, <1, 1>)  
(we, <1>)  
(tomorrow, <1>)  
(you, <1, 1, 1>)
```

### Node 1

```
(am, 3)  "am" should be "2"  
(as, 1)  
(not, 1)  
(good, 2)  
(are, 3)
```

### Node 3

```
(better, 1)  
(be, 1)  
(about, 1)  
(too, 1)  
(all, 1)  
(how, 2)  
(hi, 1)
```

### Node 2

```
(fine, 2)  
(doing, 1)  
(well, 2)  
(dad, 1)  
(i, 2)
```

### Node 4

```
(mom, 2)  
(we, 1)  
(tomorrow, 1)  
(you, 3)
```

# MapReduce Is a Compute Paradigm

- It is not a data storage paradigm
  - ▷ But any MapReduce system
  - ▷ Must read/write data from some storage system
- So MapReduce strongly linked with the idea of a distributed file system (DFS)
  - ▷ Allows data to be stored/accessed across machines in a network
  - ▷ Abstracts away differences between local and remote data
  - ▷ Same API to read/write data
  - ▷ No matter where data is located in the network

# Distributed File Systems for MR

- DFSs have been around for a long time
    - ▷ First widely used DFS was Sun's NFS, first introduced in 1985
  - Unlike classical DFS...
    - ▷ MapReduce DFS sits on top of each machine's OS
    - ▷ Lives in “user space”
    - ▷ The OS is not aware of the DFS
    - ▷ You can't mount it anywhere
  - Why on top of, not in the OS?
    - ▷ Heterogeneity no problem
    - ▷ Easily portable (JVM)
- MapReduce / Hadoop getting less popular, but the DFS is popular
- Even as MapReduce becomes less popular, MR DFS lives on!

# MapReduce vs. HPC

- MapReduce pros
  - ▷ MUCH lower programmer burden than HPC
  - ▷ No synchronization, parallelism implicit
  - ▷ Data and task placement automatic
  - ▷ Built-in fault tolerance
  - ▷ Works with (almost!) arbitrarily-sized data
  - ▷ Out-of-core execution is no problem

# MapReduce vs. HPC

- MapReduce cons
  - ▷ Standard softwares are JVM-based
  - ▷ Not suitable for communication-heavy tasks...
  - ▷ ...only communication is via the shuffle
  - ▷ Assumes BIG data... always reads/writes data from DFS

# Questions?

# Big Data Two: Beyond MapReduce

Chris Jermaine

Rice University



# Most Popular “Pure” MapReduce Software

- Is called Hadoop
  - ▷ Runs on JVM (like most Big Data software)
  - ▷ Includes MapReduce functionality
  - ▷ Plus the Hadoop distributed file system (HDFS)
- Hadoop popularity peaked around 2015...
- Has been declining since then
- Why? Were several issues...

# Hadoop MR Word Count Java Code

```
import java.util.*;

import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.IntWritable;

public class WordCount {

    public static int main(String[] args) throws Exception {

        // if we got the wrong number of args, then exit
        if (args.length != 4 || !args[0].equals ("-r")) {
            System.out.println("usage: WordCount -r <num reducers> <input> <output>");
            return -1;
        }

        // Get the default configuration object
        Configuration conf = new Configuration ();

        // now create the MapReduce job
        Job job = new Job (conf);
        job.setJobName ("WordCount");

        // we'll output text/int pairs (since we have words as keys and counts as values)
        job.setMapOutputKeyClass (Text.class);
        job.setMapOutputValueClass (IntWritable.class);

        // again we'll output text/int pairs (since we have words as keys and counts as values)
        job.setOutputKeyClass (Text.class);
        job.setOutputValueClass (IntWritable.class);

        // tell Hadoop the mapper and the reducer to use
        job.setMapperClass (WordCountMapper.class);
        job.setCombinerClass (WordCountReducer.class);
        job.setReducerClass (WordCountReducer.class);

        // we'll be reading in a text file, so we can use Hadoop's built-in TextInputFormat
        job.setInputFormatClass (TextInputFormat.class);

        // we can use Hadoop's built-in TextOutputFormat for writing out the output text file

        // set the input and output paths
        TextInputFormat.setInputPaths (job, args[2]);
        TextOutputFormat.setOutputPath (job, new Path (args[3]));

        // set the number of reduce paths
        try {
            job.setNumReduceTasks (Integer.parseInt (args[1]));
        } catch (Exception e) {
            System.out.println("usage: WordCount -r <num reducers> <input> <output>");
            return -1;
        }

        // force the mappers to handle one megabyte of input data each
        TextInputFormat.setMinInputSplitSize (job, 1024 * 1024);
        TextInputFormat.setMaxInputSplitSize (job, 1024 * 1024);

        // this tells Hadoop to ship around the jar file containing "WordCount.class" to all of
        // the different
        // nodes so that they can run the job
        job.setJarByClass (WordCount.class);

        // submit the job and wait for it to complete!
        int exitCode = job.waitForCompletion (true) ? 0 : 1;
        return exitCode;
    }
}
```

Not pretty!  
Programmer burden too high...

# Many Felt MapReduce Too Slow

- Data reread from DFS for each MR job
- Bad for iterative data processing
  - ▷ Example: most machine learning uses gradient descent
  - ▷ Need to make 100's or 1000's of passes over data
  - ▷ Re-evaluating gradient at various points

# MR API is Too Restrictive

- Can only do Map
- Or MapReduce
- Everything else in terms of those operations
- Unless you cheat
  - ▷ Ex: Mappers/Reducers talk to each other using sockets
  - ▷ But then why not just go with C/MPI?

# Result: MapReduce Used Less and Less

- Are now an entire ecosystem of alternative softwares
  - ▷ Spark, Flink, etc.
- For use in both streaming and batch processing applications
- Generally oriented more towards in-memory computing
- Have far more expressive APIs
- We will focus on Spark

# Apache Spark

- #1 Hadoop MapReduce killer
- What is Spark?
  - ▷ Platform for efficient distributed data analytics
- Runs on the JVM
- Written in Scala
  - ▷ But has bindings for Java, Scala, Python, R
  - ▷ Python nice for data analytics (NumPy, SciPy)... will focus there
- Doesn't do storage
  - ▷ Focus exclusively on compute
  - ▷ Commonly used with HDFS, S3, HBase, etc.

# RDDs

- Basic abstraction: Resilient Distributed Data Set (RDD)
- RDD is a data set buffered in RAM by Spark
  - ▷ Distributed across machines in cluster
  - ▷ To create and load an RDD (in Python shell):

```
myRDD = sc.textFile (someFileName) # sc is the Spark context  
# or else...  
data = [1, 2, 3, 4, 5]  
myRDD = sc.parallelize (data) # or  
myRDD = sc.parallelize (range (20000)) # or...
```

# Computations: Series of Xforms Over RDDs

- Example: word count

```
def countWords (fileName):  
    textFile = sc.textFile (fileName)  
    lines = textFile.flatMap (lambda line: line.split("_"))  
    counts = lines.map (lambda word: (word, 1))  
    aggCounts = counts.reduceByKey (lambda a, b: a + b)  
    return aggCounts.top (200, key=lambda p: p[1])
```

- What transforms do we see here?
  - ▷ flatMap, map, reduceByKey, top
- Let's go through them
- But first, quick review of lambdas...

▷ Fundamental to programming in Spark



# What's a Lambda?

important to Spark

- Basically, a function that that we can pass like a variable
- Key ability: can “capture” its surroundings at creation

```
def addTwelveToResult (myLambda) :  
    return myLambda (3) + 12
```

```
a = 23
```

```
aCoolLambda = lambda x : x + a
```

```
addTwelveToResult (aCoolLambda) # prints 38
```

```
a = 45
```

```
addTwelveToResult (aCoolLambda) # prints ???
```

# Lambdas and Comprehensions

- Lambdas can return many items

```
def sumThem (myLambda):  
    tot = 0  
    for a in myLambda ():  
        tot = tot + a  
    return tot  
  
x = np.array([1, 2, 3, 4, 5])  
iter = lambda : (j for j in x)  
sumThem (iter) # prints 15
```

# Spark Operation: flatMap ()

```
def countWords (fileName):  
    textFile = sc.textFile (fileName)  
    lines = textFile.flatMap (lambda line: line.split("_"))
```

split every time you see a space

- Process every data item in the RDD
- Apply lambda to it
- Lambda argument to return zero or more results
- Each result added into resulting RDD

creates a new RDD list of words

# Spark Operation: map ()

```
def countWords (fileName):  
    textFile = sc.textFile (fileName)  
    lines = textFile.flatMap (lambda line: line.split("_"))  
    counts = lines.map (lambda word: (word, 1))
```

- Process every data item in the RDD
- Apply lambda to it
- The lambda must return exactly one result

could have used flatmap instead of map, but there is some performance benefit to using map which states there is 1 result (vs. zero or more)

map expects to produce single output = list?  
flatMap provide a list and create a new RDD?

# Spark Operation: reduceByKey ()

```
def countWords (fileName):  
    textFile = sc.textFile (fileName)  
    lines = textFile.flatMap (lambda line: line.split("_"))  
    counts = lines.map (lambda word: (word, 1))  
    aggCounts = counts.reduceByKey (lambda a, b: a + b)
```

- Data must be  $(Key, Value)$  pairs
- Shuffle so that all  $(K, V)$  pairs with same  $K$  on same machine
- Organize into  $(K, (V_1, V_2, ..., V_n))$  pairs
- Use the lambda to “reduce” the list to a single value

# Spark Operation: top ()

```
def countWords (fileName):  
    textFile = sc.textFile (fileName)  
    lines = textFile.flatMap (lambda line: line.split("_"))  
    counts = lines.map (lambda word: (word, 1))  
    aggCounts = counts.reduceByKey (lambda a, b: a + b)  
    retrun aggCounts.top (200, key=lamda p: p[1])
```

- Data must be (*Key, Value*) pairs
- Takes two params... first is number to return
- Second (optional): lambda to use to obtain key for comparison
- Note: top collects an RDD, moving from cloud to local
- So result is not an RDD

top is a collection operation

top returns some sort of a list data structure

# An Important Note

- Spark uses lazy evaluation...
- If I run this code:

```
textFile = sc.textFile (fileName)
lines = textFile.flatMap (lambda line: line.split("_"))
counts = lines.map (lambda word: (word, 1))
aggCounts = counts.reduceByKey (lambda a, b: a + b)
```

- Nothing happens! (Other than Spark remembers the ops)
  - ▷ Spark does not execute until an attempt made to collect an RDD
  - ▷ When we hit `top()`, then all of these are executed
- Why do this?
  - ▷ By waiting until last possible second, opportunities for “pipelining” exploited
  - ▷ Only ops that require a shuffle can’t be pipelined

# Some Other, More Advanced Ops

- `groupByKey ()`, `join ()`, `reduce ()`, `aggregate ()`



## groupByKey ()

- Data must be  $(Key, Value)$  pairs
- Shuffle so that all  $(K, V)$  pairs with same  $K$  on same machine
- Organize into  $(K, \langle V_1, V_2, \dots, V_n \rangle)$  pairs
- Store each list as a `ResultIterable` for future processing
- Like `reduceByKey ()` but without the reduce

## join ()

- Given two data sets *rddOne*, *rddTwo* of (*Key*, *Value*) pairs
- We join them using:

`rddOne.join (rddTwo)`

- Returns (*K*, (*V*<sub>1</sub>, *V*<sub>2</sub>)) pairs
- Constructed from all (*K*<sub>1</sub>, *V*<sub>1</sub>) from *rddOne*, (*K*<sub>2</sub>, *V*<sub>2</sub>) from *rddTwo*, where *K*<sub>1</sub> = *K*<sub>2</sub>

# join () (Continued)

- Example:
  - ▷ *rddOne* is  $\{(red, 9), (blue, 7), (red, 12), (green, 4)\}$
  - ▷ *rddTwo* is  $\{(blue, up), (green, down), (green, behind)\}$
  - ▷ Result of join is  $\{(blue, (7, up)), (green, (4, down)), (green, (4, behind))\}$
- Can blow up RDD size if join is many-to-many
- Requires expensive shuffle!

# Why Do We Join?

- Allows “communication”
- Imagine fluid simulation
  - ▷ State stored in RDD of  $(gridCell, (xPos, yPos, state))$  pairs
- Join this RDD with itself
  - ▷ To get  $(gridCell, ((xPos_1, yPos_1, state_1), (xPos_2, yPos_2, state_2)))$  pairs
  - ▷ Then map  $()$  to get  $((xPos_1, yPos_1, state_1), (xPos_2, yPos_2, state_2))$  pairs
  - ▷ groupByKey  $()$  to get  $(xPos_1, yPos_1, state_1)$  with all states in same grid cell
  - ▷ Finally, evoke a lambda to update  $(xPos_1, yPos_1, state_1)$  based on neighbors

# reduce ()

- Unlike past operations, this is not a transform from RDD to RDD
  - ▷ This is like `top ()`, moves result back to Python
- Repeatedly applies a lambda to each item in RDD to get single result

```
>>> myData = sc.parallelize (range(20000))
>>> myData.reduce (lambda a, b: a + b)
199990000
```

# aggregate ()

- With `reduce ()` you aggregate directly, can be restrictive...
- Example: RDD is  $\{(red, 9), (blue, 7), (red, 12), (green, 4)\}$ 
  - ▷ Want dictionary where value is sum for each unique color
  - ▷ Cannot use `reduce ()`
  - ▷ It only “sums” up the items in the input RDD directly
  - ▷ Two inputs and output must be the same type
  - ▷ How do I get the desired output dictionary by defining  $+$  in  $((red, 9) + (blue, 7)) + (red, 12) + (green, 4)$ ?

## aggregate () (Continued)

- Data must be  $(Key, Value)$  pairs
- Organize into  $(K, \langle V_1, V_2, \dots, V_n \rangle)$  pairs
- Then aggregate the list, like `reduce ()`
- `aggregate ()` takes three args
  - ▷ The “zero” to init the aggregation
  - ▷ Lambda that takes  $X_1, X_2$  and aggs them, where  $X_1$  already aggregated,  $X_2$  not
  - ▷ Lambda that takes  $X_1, X_2$  and aggs them, where both aggregated

## aggregate () (Example)

```
def add (dict, tuple):  
    result = {}  
    for key in dict:  
        result[key] = dict[key]  
    if (tuple[0] in result):  
        result[tuple[0]] += tuple[1]  
    else:  
        result[tuple[0]] = tuple[1]  
    return result
```



## aggregate () (Example)

```
def combine (dict1, dict2):  
    result = {}  
    for key in dict1:  
        result[key] = dict1[key]  
    for key in dict2:  
        if (key in result):  
            result[key] += dict2[key]  
        else:  
            result[key] = dict2[key]  
    return result
```

## aggregate () (Example)

```
>>> myRdd = sc.parallelize ([('red', 9), ('blue', 7),  
... ('red', 12), ('green', 4)])  
>>> myRdd.aggregate ({}, lambda x, y: add (x, y),  
... lambda x, y: combine (x, y))  
  
{'blue': 7, 'green': 4, 'red': 21}
```

# Closing Thoughts

- When is Spark/MapReduce a better option than HPC?
  - ▷ When your pipeline is heavily data-oriented
  - ▷ Or when your compute is (relatively) loosely coupled
- Key benefits compared to HPC
  - ▷ Built in fault tolerance
  - ▷ Better support for BIG data
  - ▷ Much higher programmer productivity
- Will continue to take market share from HPC
  - ▷ You see academic papers with both MPI, Spark implementations
  - ▷ But not everything can move to Spark

# Questions?