

Download data

Data available via scikit-learn

<https://scikit-learn.org/stable/modules/classes.html#module-sklearn.datasets> (<https://scikit-learn.org/stable/modules/classes.html#module-sklearn.datasets>)

```
In [14]: from sklearn import datasets

##### classifiction (supervised/descrete) #####
wine_bunch = datasets.load_wine()
print(type(wine_bunch))
print(wine_bunch.keys())
# print(wine_bunch.DESCR)

##### regression (supervised/continuous) #####
reg_bunch = datasets.load_boston()
print(reg_bunch.keys())

##### clustering (unsupervised/descrete) #####
#radmon generated data
# gene expression
##### dimensionality reduction (unuspervised/continuous) #####
#https://idyll.pub/post/dimensionality-reduction-293e465c2a3443e8941b016d/
#gene expression
```

```
<class 'sklearn.utils.Bunch'>
dict_keys(['data', 'target', 'target_names', 'DESCR', 'feature_names'])
.. _wine_dataset:
```

Wine recognition dataset

****Data Set Characteristics:****

```
:Number of Instances: 178 (50 in each of three classes)
:Number of Attributes: 13 numeric, predictive attributes and the class
:Attribute Information:
```

- Alcohol
- Malic acid
- Ash
- Alcalinity of ash
- Magnesium
- Total phenols
- Flavanoids
- Nonflavanoid phenols
- Proanthocyanins
- Color intensity
- Hue
- OD280/OD315 of diluted wines
- Proline

```
- class:
```

- class_0
- class_1
- class_2

```
:Summary Statistics:
```

	Min	Max	Mean	SD
Alcohol:	11.0	14.8	13.0	0.8
Malic Acid:	0.74	5.80	2.34	1.12
Ash:	1.36	3.23	2.36	0.27
Alcalinity of Ash:	10.6	30.0	19.5	3.3
Magnesium:	70.0	162.0	99.7	14.3
Total Phenols:	0.98	3.88	2.29	0.63
Flavanoids:	0.34	5.08	2.03	1.00
Nonflavanoid Phenols:	0.13	0.66	0.36	0.12
Proanthocyanins:	0.41	3.58	1.59	0.57
Colour Intensity:	1.3	13.0	5.1	2.3
Hue:	0.48	1.71	0.96	0.23
OD280/OD315 of diluted wines:	1.27	4.00	2.61	0.71
Proline:	278	1680	746	315

```
:Missing Attribute Values: None
```

```
:Class Distribution: class_0 (59), class_1 (71), class_2 (48)
```

```
:Creator: R.A. Fisher
```

```
:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)
```

```
:Date: July, 1988
```

This is a copy of UCI ML Wine recognition datasets.

<https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data>

The data is the results of a chemical analysis of wines grown in the same region in Italy by three different cultivators. There are thirteen different measurements taken for different constituents found in the three types of wine.

Original Owners:

Forina, M. et al, PARVUS -

An Extendible Package for Data Exploration, Classification and Correlation.
Institute of Pharmaceutical and Food Analysis and Technologies,
Via Brigata Salerno, 16147 Genoa, Italy.

Citation:

Lichman, M. (2013). UCI Machine Learning Repository
[<https://archive.ics.uci.edu/ml>]. Irvine, CA: University of California,
School of Information and Computer Science.

.. topic:: References

(1) S. Aeberhard, D. Coomans and O. de Vel,
Comparison of Classifiers in High Dimensional Settings,
Tech. Rep. no. 92-02, (1992), Dept. of Computer Science and Dept. of
Mathematics and Statistics, James Cook University of North Queensland.
(Also submitted to Technometrics).

The data was used with many others for comparing various
classifiers. The classes are separable, though only RDA
has achieved 100% correct classification.
(RDA : 100%, QDA 99.4%, LDA 98.9%, 1NN 96.1% (z-transformed data))
(All results using the leave-one-out technique)

(2) S. Aeberhard, D. Coomans and O. de Vel,
"THE CLASSIFICATION PERFORMANCE OF RDA"
Tech. Rep. no. 92-01, (1992), Dept. of Computer Science and Dept. of
Mathematics and Statistics, James Cook University of North Queensland.
(Also submitted to Journal of Chemometrics).

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

Pandas overview

```
In [16]: import pandas as pd
import numpy as np
wine = pd.DataFrame(wine_bunch.data, columns=wine_bunch.feature_names)
display(wine)

# # play with these functions #
# print(list(wine))
# print(wine.proline)
# print(wine.iloc[10:20,])
# # print(wine.iloc[10:20,['ash','magnesium']]) # error bec iloc but col name
instead of index
# print(wine.iloc[10:20,0:3])
# print(wine.loc[10:20,['ash','magnesium']])
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflava
0	14.23	1.71	2.43	15.6	127.0	2.80	3.06	
1	13.20	1.78	2.14	11.2	100.0	2.65	2.76	
2	13.16	2.36	2.67	18.6	101.0	2.80	3.24	
3	14.37	1.95	2.50	16.8	113.0	3.85	3.49	
4	13.24	2.59	2.87	21.0	118.0	2.80	2.69	
5	14.20	1.76	2.45	15.2	112.0	3.27	3.39	
6	14.39	1.87	2.45	14.6	96.0	2.50	2.52	
7	14.06	2.15	2.61	17.6	121.0	2.60	2.51	
8	14.83	1.64	2.17	14.0	97.0	2.80	2.98	
9	13.86	1.35	2.27	16.0	98.0	2.98	3.15	
10	14.10	2.16	2.30	18.0	105.0	2.95	3.32	
11	14.12	1.48	2.32	16.8	95.0	2.20	2.43	
12	13.75	1.73	2.41	16.0	89.0	2.60	2.76	
13	14.75	1.73	2.39	11.4	91.0	3.10	3.69	
14	14.38	1.87	2.38	12.0	102.0	3.30	3.64	
15	13.63	1.81	2.70	17.2	112.0	2.85	2.91	
16	14.30	1.92	2.72	20.0	120.0	2.80	3.14	
17	13.83	1.57	2.62	20.0	115.0	2.95	3.40	
18	14.19	1.59	2.48	16.5	108.0	3.30	3.93	
19	13.64	3.10	2.56	15.2	116.0	2.70	3.03	
20	14.06	1.63	2.28	16.0	126.0	3.00	3.17	
21	12.93	3.80	2.65	18.6	102.0	2.41	2.41	
22	13.71	1.86	2.36	16.6	101.0	2.61	2.88	
23	12.85	1.60	2.52	17.8	95.0	2.48	2.37	
24	13.50	1.81	2.61	20.0	96.0	2.53	2.61	
25	13.05	2.05	3.22	25.0	124.0	2.63	2.68	
26	13.39	1.77	2.62	16.1	93.0	2.85	2.94	
27	13.30	1.72	2.14	17.0	94.0	2.40	2.19	
28	13.87	1.90	2.80	19.4	107.0	2.95	2.97	
29	14.02	1.68	2.21	16.0	96.0	2.65	2.33	
...	
148	13.32	3.24	2.38	21.5	92.0	1.93	0.76	
149	13.08	3.90	2.36	21.5	113.0	1.41	1.39	
150	13.50	3.12	2.62	24.0	123.0	1.40	1.57	
151	12.79	2.67	2.48	22.0	112.0	1.48	1.36	

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflava
152	13.11	1.90	2.75	25.5	116.0	2.20	1.28	
153	13.23	3.30	2.28	18.5	98.0	1.80	0.83	
154	12.58	1.29	2.10	20.0	103.0	1.48	0.58	
155	13.17	5.19	2.32	22.0	93.0	1.74	0.63	
156	13.84	4.12	2.38	19.5	89.0	1.80	0.83	
157	12.45	3.03	2.64	27.0	97.0	1.90	0.58	
158	14.34	1.68	2.70	25.0	98.0	2.80	1.31	
159	13.48	1.67	2.64	22.5	89.0	2.60	1.10	
160	12.36	3.83	2.38	21.0	88.0	2.30	0.92	
161	13.69	3.26	2.54	20.0	107.0	1.83	0.56	
162	12.85	3.27	2.58	22.0	106.0	1.65	0.60	
163	12.96	3.45	2.35	18.5	106.0	1.39	0.70	
164	13.78	2.76	2.30	22.0	90.0	1.35	0.68	
165	13.73	4.36	2.26	22.5	88.0	1.28	0.47	
166	13.45	3.70	2.60	23.0	111.0	1.70	0.92	
167	12.82	3.37	2.30	19.5	88.0	1.48	0.66	
168	13.58	2.58	2.69	24.5	105.0	1.55	0.84	
169	13.40	4.60	2.86	25.0	112.0	1.98	0.96	
170	12.20	3.03	2.32	19.0	96.0	1.25	0.49	
171	12.77	2.39	2.28	19.5	86.0	1.39	0.51	
172	14.16	2.51	2.48	20.0	91.0	1.68	0.70	
173	13.71	5.65	2.45	20.5	95.0	1.68	0.61	
174	13.40	3.91	2.48	23.0	102.0	1.80	0.75	
175	13.27	4.28	2.26	20.0	120.0	1.59	0.69	
176	13.17	2.59	2.37	20.0	120.0	1.65	0.68	
177	14.13	4.10	2.74	24.5	96.0	2.05	0.76	

178 rows × 13 columns



```
['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phe  
nols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intens  
ity', 'hue', 'od280/od315_of_diluted_wines', 'proline']  
0      1065.0  
1      1050.0  
2      1185.0  
3      1480.0  
4       735.0  
5      1450.0  
6      1290.0  
7      1295.0  
8      1045.0  
9      1045.0  
10     1510.0  
11     1280.0  
12     1320.0  
13     1150.0  
14     1547.0  
15     1310.0  
16     1280.0  
17     1130.0  
18     1680.0  
19      845.0  
20      780.0  
21      770.0  
22     1035.0  
23     1015.0  
24      845.0  
25      830.0  
26     1195.0  
27     1285.0  
28      915.0  
29     1035.0  
  
...  
148     650.0  
149     550.0  
150     500.0  
151     480.0  
152     425.0  
153     675.0  
154     640.0  
155     725.0  
156     480.0  
157     880.0  
158     660.0  
159     620.0  
160     520.0  
161     680.0  
162     570.0  
163     675.0  
164     615.0  
165     520.0  
166     695.0  
167     685.0  
168     750.0  
169     630.0  
170     510.0
```


171 470.0
 172 660.0
 173 740.0
 174 750.0
 175 835.0
 176 840.0
 177 560.0

Name: proline, Length: 178, dtype: float64

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	\
10	14.10	2.16	2.30	18.0	105.0	2.95	
11	14.12	1.48	2.32	16.8	95.0	2.20	
12	13.75	1.73	2.41	16.0	89.0	2.60	
13	14.75	1.73	2.39	11.4	91.0	3.10	
14	14.38	1.87	2.38	12.0	102.0	3.30	
15	13.63	1.81	2.70	17.2	112.0	2.85	
16	14.30	1.92	2.72	20.0	120.0	2.80	
17	13.83	1.57	2.62	20.0	115.0	2.95	
18	14.19	1.59	2.48	16.5	108.0	3.30	
19	13.64	3.10	2.56	15.2	116.0	2.70	

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
10	3.32		0.22	2.38	5.75	1.25
11	2.43		0.26	1.57	5.00	1.17
12	2.76		0.29	1.81	5.60	1.15
13	3.69		0.43	2.81	5.40	1.25
14	3.64		0.29	2.96	7.50	1.20
15	2.91		0.30	1.46	7.30	1.28
16	3.14		0.33	1.97	6.20	1.07
17	3.40		0.40	1.72	6.60	1.13
18	3.93		0.32	1.86	8.70	1.23
19	3.03		0.17	1.66	5.10	0.96

	od280/od315_of_diluted_wines	proline
10	3.17	1510.0
11	2.82	1280.0
12	2.90	1320.0
13	2.73	1150.0
14	3.00	1547.0
15	2.88	1310.0
16	2.65	1280.0
17	2.57	1130.0
18	2.82	1680.0
19	3.36	845.0

	alcohol	malic_acid	ash
10	14.10	2.16	2.30
11	14.12	1.48	2.32
12	13.75	1.73	2.41
13	14.75	1.73	2.39
14	14.38	1.87	2.38
15	13.63	1.81	2.70
16	14.30	1.92	2.72
17	13.83	1.57	2.62
18	14.19	1.59	2.48
19	13.64	3.10	2.56

	ash	magnesium
10	2.30	105.0

```

11  2.32      95.0
12  2.41      89.0
13  2.39      91.0
14  2.38     102.0
15  2.70     112.0
16  2.72     120.0
17  2.62     115.0
18  2.48     108.0
19  2.56     116.0
20  2.28     126.0

```

wineX is Pandas DataFrame (because we made it a DF earlier)

wineY is ndarray

```

In [19]: wineX = wine
        wineY = wine_bunch.target

        print(wineX.shape, wineY.shape)
        print(type(wineX))
        print(type(wineY))

(178, 13) (178,)
<class 'pandas.core.frame.DataFrame'>
<class 'numpy.ndarray'>

```

Supervised classification

Test / train split

```

In [32]: # random sampling
        from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(wineX, wineY, test_size=0.2, random_state=42)
        print(X_train.shape, y_train.shape)
        print(X_test.shape, y_test.shape)
        print(type(X_train))
        print(type(X_test))
        print(type(y_train))
        print(type(y_test))

(142, 13) (142,)
(36, 13) (36,)
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.frame.DataFrame'>
<class 'numpy.ndarray'>
<class 'numpy.ndarray'>

```

Repeat the split, but in the context of K-fold validation

```
In [33]: # k-fold
from sklearn.model_selection import KFold
kf = KFold(n_splits=5, random_state=42, shuffle=True)
kf.get_n_splits(wineX)
print(kf)
for train_index, test_index in kf.split(wineX):
    X_train, X_test = wineX.iloc[train_index], wineX.iloc[test_index]
    y_train, y_test = wineY[train_index], wineY[test_index]
    print(X_train.shape, y_train.shape)
    print(X_test.shape, y_test.shape)
# NOTE: wineX is a dataframe, wineY is a ndarray (syntax is different line 7 and 8)
print(type(X_train))
print(type(X_test))
print(type(y_train))
print(type(y_test))


KFold(n_splits=5, random_state=42, shuffle=True)
(142, 13) (142,)
(36, 13) (36,)
(142, 13) (142,)
(36, 13) (36,)
(142, 13) (142,)
(36, 13) (36,)
(143, 13) (143,)
(35, 13) (35,)
(143, 13) (143,)
(35, 13) (35,)
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.frame.DataFrame'>
<class 'numpy.ndarray'>
<class 'numpy.ndarray'>
```

KNN classification

Visualize the data

```
In [34]: %matplotlib inline
# K Nearest Neighbors
import matplotlib.pyplot as plt
print(list(X_train))
plt.scatter(X_train['alcalinity_of_ash'], X_train['total_phenols'], c=y_train)
plt.show()
# CHANGE which features that you are plotting
# Can only visualize on 2 dimensions but algorithm will cluster on all dimensions
```

```
['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']
```



```
In [30]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=3, leaf_size=10)
neigh.fit(X_train, y_train)
print(neigh)
# print(neigh.predict(X_test))
# print(neigh.predict_proba(X_test))
# # TEST SCORE
# print(neigh.score(X_test,y_test))
# # TRAINING SCORE
# print(neigh.score(X_train,y_train))
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=10, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=3, p=2,
                    weights='uniform')
```



```
[0 0 0 2 0 0 0 0 0 0 0 1 1 0 1 1 1 1 1 1 1 1 1 2 2 2 1 0 0 2 2 0 2]
```

```
[[1.      0.      0.      ]
 [1.      0.      0.      ]
 [0.66666667 0.33333333 0.      ]
 [0.      0.33333333 0.66666667]
 [1.      0.      0.      ]
 [1.      0.      0.      ]
 [1.      0.      0.      ]
 [1.      0.      0.      ]
 [1.      0.      0.      ]
 [1.      0.      0.      ]
 [0.      0.66666667 0.33333333]
 [0.      0.66666667 0.33333333]
 [0.66666667 0.33333333 0.      ]
 [0.      0.66666667 0.33333333]
 [0.      0.66666667 0.33333333]
 [0.      0.66666667 0.33333333]
 [0.      1.      0.      ]
 [0.      0.66666667 0.33333333]
 [0.      1.      0.      ]
 [0.      0.66666667 0.33333333]
 [0.      1.      0.      ]
 [0.      0.66666667 0.33333333]
 [0.      1.      0.      ]
 [0.      0.66666667 0.33333333]
 [0.      0.33333333 0.66666667]
 [0.      0.33333333 0.66666667]
 [0.      0.33333333 0.66666667]
 [0.      0.33333333 0.66666667]
 [0.      0.66666667 0.33333333]
 [0.66666667 0.33333333 0.      ]
 [0.33333333 0.33333333 0.33333333]
 [0.      0.33333333 0.66666667]
 [0.      0.33333333 0.66666667]
 [0.66666667 0.      0.33333333]
 [0.      0.33333333 0.66666667]]
0.8285714285714286
0.8671328671328671
```

```
In [35]: print(list(X_train))
y_pred = neigh.predict(X_test)

#PLOT RESULTS
plt.title("Original Y")
plt.scatter(X_test['alcalinity_of_ash'], X_test['total_phenols'], c=y_test)
plt.show()
plt.title("Predicted Y")
plt.scatter(X_test['alcalinity_of_ash'], X_test['total_phenols'], c=y_pred)
plt.show()
```

['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']



Visualizing the results

TRY ON DIFFERENT SPLIT!!

```
In [39]: # run cell below plotting function heading first before running this cell (out of order)

from sklearn.metrics import confusion_matrix
from sklearn.utils.multiclass import unique_labels

class_names = wine_bunch.target_names

plot_confusion_matrix(y_test, y_pred, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plot_confusion_matrix(y_test, y_pred, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')
```



Confusion matrix, without normalization

```
[[10  0  1]
 [ 1 13  0]
 [ 3  1  6]]
```

Normalized confusion matrix

```
[[0.90909091 0.         0.09090909]
 [0.07142857 0.92857143 0.         ]
 [0.3         0.1         0.6         ]]
```

Out[39]: <matplotlib.axes._subplots.AxesSubplot at 0x1c2e629aa20>



Plotting function

Reference:

https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html#sphx-glr-auto-examples-model-selection-plot-confusion-matrix-py (https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html#sphx-glr-auto-examples-model-selection-plot-confusion-matrix-py)


```

In [38]: def plot_confusion_matrix(y_true, y_pred, classes,
                                   normalize=False,
                                   title=None,
                                   cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if not title:
        if normalize:
            title = 'Normalized confusion matrix'
        else:
            title = 'Confusion matrix, without normalization'

    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    # Only use the labels that appear in the data
    classes = classes[unique_labels(y_true, y_pred)]
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    fig, ax = plt.subplots()
    im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
    ax.figure.colorbar(im, ax=ax)
    # We want to show all ticks...
    ax.set(xticks=np.arange(cm.shape[1]),
           yticks=np.arange(cm.shape[0]),
           # ... and label them with the respective list entries
           xticklabels=classes, yticklabels=classes,
           title=title,
           ylabel='True label',
           xlabel='Predicted label')

    # Rotate the tick labels and set their alignment.
    plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
              rotation_mode="anchor")

    # Loop over data dimensions and create text annotations.
    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            ax.text(j, i, format(cm[i, j], fmt),
                    ha="center", va="center",
                    color="white" if cm[i, j] > thresh else "black")
    fig.tight_layout()
    return ax

```

Excercise 1 (30minutes)

```
In [61]: # 1. Perform K-fold KNN cross validation on the wine data
# 2. Modify KNN hyper parameters and see how it effects the results

# #pseudo code:
# create k-folds
# for each split:
#     train KNN
#     predict for the split
#     see the score

# k-fold
from sklearn.model_selection import KFold
kf = KFold(n_splits=6,random_state=42, shuffle=True)
kf.get_n_splits(wineX)
print(kf)

for train_index, test_index in kf.split(wineX):
    X_train, X_test = wineX.iloc[train_index], wineX.iloc[test_index]
    y_train, y_test = wineY[train_index], wineY[test_index]
    print(X_train.shape, y_train.shape)
    print(X_test.shape, y_test.shape)
    neigh = KNeighborsClassifier(n_neighbors=3, leaf_size=10)
    neigh.fit(X_train, y_train)
    # print(neigh)
    # print(neigh.predict(X_test))
    # print(neigh.predict_proba(X_test))
    # TEST SCORE
    print("TEST score {:.f}".format(neigh.score(X_test,y_test)))
    # TRAINING SCORE
    print("TRAINING score {:.f}".format(neigh.score(X_train,y_train)))
```

```
KFold(n_splits=6, random_state=42, shuffle=True)
```

```
(148, 13) (148,)
```

```
(30, 13) (30,)
```

```
TEST score 0.766667
```

```
TRAINING score 0.864865
```

```
(148, 13) (148,)
```

```
(30, 13) (30,)
```

```
TEST score 0.766667
```

```
TRAINING score 0.858108
```

```
(148, 13) (148,)
```

```
(30, 13) (30,)
```

```
TEST score 0.600000
```

```
TRAINING score 0.864865
```

```
(148, 13) (148,)
```

```
(30, 13) (30,)
```

```
TEST score 0.700000
```

```
TRAINING score 0.864865
```

```
(149, 13) (149,)
```

```
(29, 13) (29,)
```

```
TEST score 0.551724
```

```
TRAINING score 0.872483
```

```
(149, 13) (149,)
```

```
(29, 13) (29,)
```

```
TEST score 0.827586
```

```
TRAINING score 0.879195
```

Supervised regression

Test / train split

```
In [72]: ##### regression (supervised/continuous) #####
reg_bunch = datasets.load_boston()
print(reg_bunch.keys())
print(reg_bunch.data)

# random sampling
from sklearn import preprocessing
reg_bunch.data = preprocessing.scale(reg_bunch.data)
print(reg_bunch.DESCR)
regX = pd.DataFrame(reg_bunch.data, columns=reg_bunch.feature_names)
regY = reg_bunch.target
print("regX sample: \n", regX.iloc[0:10])
print("regY sample: \n", regY[0:10])

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(regX, regY, test_size=0.2,
random_state=42)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers tha

t address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

regX sample:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE \
0	-0.419782	0.284830	-1.287909	-0.272599	-0.144217	0.413672	-0.120013
1	-0.417339	-0.487722	-0.593381	-0.272599	-0.740262	0.194274	0.367166
2	-0.417342	-0.487722	-0.593381	-0.272599	-0.740262	1.282714	-0.265812
3	-0.416750	-0.487722	-1.306878	-0.272599	-0.835284	1.016303	-0.809889
4	-0.412482	-0.487722	-1.306878	-0.272599	-0.835284	1.228577	-0.511180
5	-0.417044	-0.487722	-1.306878	-0.272599	-0.835284	0.207096	-0.351157
6	-0.410243	0.048772	-0.476654	-0.272599	-0.265154	-0.388411	-0.070229
7	-0.403696	0.048772	-0.476654	-0.272599	-0.265154	-0.160466	0.978808
8	-0.395935	0.048772	-0.476654	-0.272599	-0.265154	-0.931206	1.117494
9	-0.400729	0.048772	-0.476654	-0.272599	-0.265154	-0.399808	0.616090

	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.140214	-0.982843	-0.666608	-1.459000	0.441052	-1.075562
1	0.557160	-0.867883	-0.987329	-0.303094	0.441052	-0.492439
2	0.557160	-0.867883	-0.987329	-0.303094	0.396427	-1.208727
3	1.077737	-0.752922	-1.106115	0.113032	0.416163	-1.361517
4	1.077737	-0.752922	-1.106115	0.113032	0.441052	-1.026501
5	1.077737	-0.752922	-1.106115	0.113032	0.410571	-1.043322
6	0.839244	-0.523001	-0.577519	-1.505237	0.426798	-0.031268
7	1.024638	-0.523001	-0.577519	-1.505237	0.441052	0.910700
8	1.087196	-0.523001	-0.577519	-1.505237	0.328448	2.421774
9	1.329635	-0.523001	-0.577519	-1.505237	0.329325	0.623344

regY sample:

```
[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9]
(404, 13) (404,)
(102, 13) (102,)
```

Regression

https://scikit-learn.org/stable/modules/linear_model.html (https://scikit-learn.org/stable/modules/linear_model.html)


```
In [92]: ##### Linear regression #####
from sklearn import linear_model

reg = linear_model.LinearRegression()
reg.fit(X_train, y_train)
print(reg.coef_)
y_pred = reg.predict(X_test)
print("LR Train score: {:.f}".format(reg.score(X_train,y_train)))
print("LR Test score: {:.f}".format(reg.score(X_test,y_test)))
plt.scatter(y_test, y_pred)
plt.show()
```

```
[-0.97149423  0.70155562  0.27675212  0.70653152 -1.99143043  3.11571836
 -0.17706021 -3.04577065  2.28278471 -1.79260468 -1.97995351  1.12649864
 -3.62814937]
```

LR Train score: 0.750886

LR Test score: 0.668759




```

In [93]: ##### RIDGE regression #####
from sklearn.linear_model import Ridge

def find_alpha():
    # adopted from: https://towardsdatascience.com/ridge-regression-for-better
    -usage-2f19b3a202db
    ridge_reg = Ridge(alpha=0)
    ridge_reg.fit(X_train, y_train)
    ridge_df = pd.DataFrame({'variable': reg_bunch.feature_names, 'estimate':
ridge_reg.coef_})
    ridge_train_pred = []
    ridge_test_pred = []
    # iterate lambdas


    # iterate lambdas
    for alpha in np.arange(0, 200, 1):
        # training
        ridge_reg = Ridge(alpha=alpha)
        ridge_reg.fit(X_train, y_train)
        var_name = 'estimate' + str(alpha)
        ridge_df[var_name] = ridge_reg.coef_
        # prediction
        ridge_train_pred.append(ridge_reg.predict(X_train))
        ridge_test_pred.append(ridge_reg.predict(X_test))

    # organize dataframe
    ridge_df = ridge_df.set_index('variable').T.rename_axis('estimate').rename
_axis(1).reset_index()
    # plot betas by Lambda
    fig, ax = plt.subplots(figsize=(10, 5))
    ax.plot(ridge_df.RM, 'r', ridge_df.ZN, 'g', ridge_df.RAD, 'b', ridge_df.CR
IM, 'c', ridge_df.TAX, 'y')
    ax.axhline(y=0, color='black', linestyle='--')
    ax.set_xlabel("Lambda")
    ax.set_ylabel("Beta Estimate")
    ax.set_title("Ridge Regression Trace", fontsize=16)
    ax.legend(labels=['Room', 'Residential Zone', 'Highway Access', 'Crime Rate',
'Tax'])
    ax.grid(True)
    plt.show()

# Rooms is the features that has the most significance
find_alpha()

ridge = Ridge(alpha=20.0)
ridge.fit(X_train, y_train)
print(ridge.coef_)
y_pred = ridge.predict(X_test)
print("RIDGE Train score: {:.f}".format(ridge.score(X_train,y_train)))
print("RIDGE Test score: {:.f}".format(ridge.score(X_test,y_test)))
plt.scatter(y_test, y_pred)
plt.show()


```



```
[ -0.84416652  0.46198062 -0.03418626  0.75114482 -1.53723325  3.17316115
 -0.1880602  -2.48151078  1.31849282 -0.97365431 -1.83807721  1.08177436
 -3.40377171]
```

RIDGE Train score: 0.747905

RIDGE Test score: 0.663563



```

In [94]: ##### LASSO regression #####
from sklearn import linear_model

def find_alpha():
    ridge_reg = linear_model.Lasso(alpha=0.01)
    ridge_reg.fit(X_train, y_train)
    ridge_df = pd.DataFrame({'variable': reg_bunch.feature_names, 'estimate':
    ridge_reg.coef_})
    ridge_train_pred = []
    ridge_test_pred = []
    # iterate lambdas

    # iterate lambdas
    for alpha in np.arange(0.1, 1, 0.1):
        # training
        ridge_reg = linear_model.Lasso(alpha=alpha)
        ridge_reg.fit(X_train, y_train)
        var_name = 'estimate' + str(alpha)
        ridge_df[var_name] = ridge_reg.coef_
        # prediction
        ridge_train_pred.append(ridge_reg.predict(X_train))
        ridge_test_pred.append(ridge_reg.predict(X_test))

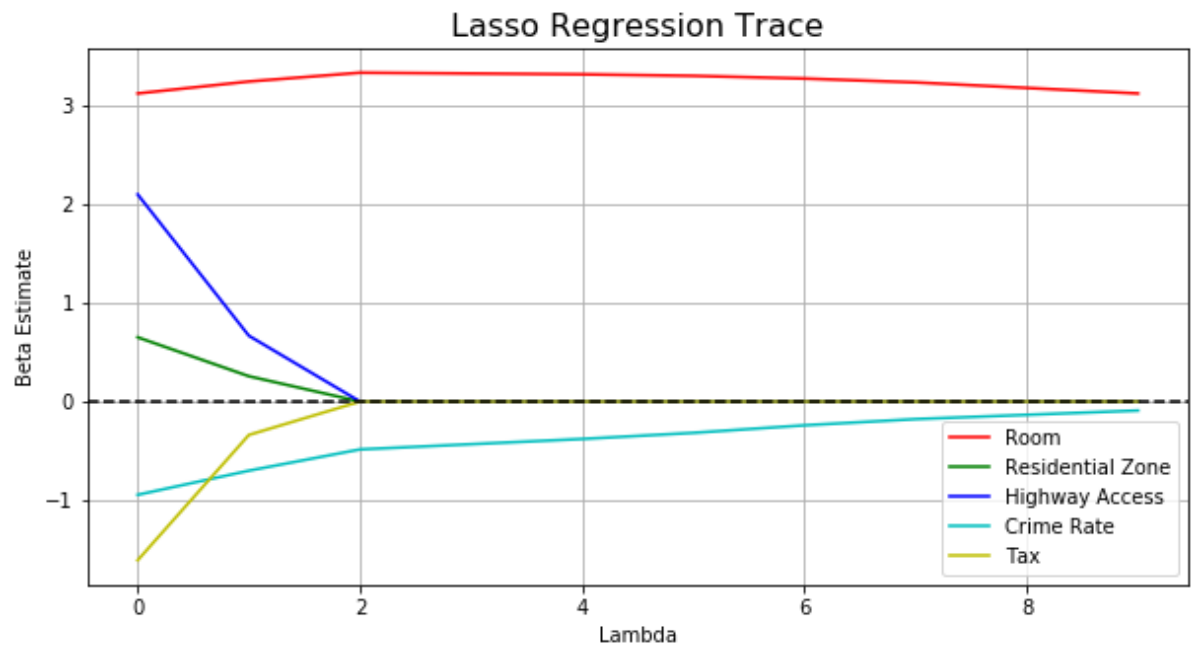
    # organize dataframe
    ridge_df = ridge_df.set_index('variable').T.rename_axis('estimate').rename
    _axis(1).reset_index()
    # plot betas by lambda
    fig, ax = plt.subplots(figsize=(10, 5))
    ax.plot(ridge_df.RM, 'r', ridge_df.ZN, 'g', ridge_df.RAD, 'b', ridge_df.CR
    IM, 'c', ridge_df.TAX, 'y')
    ax.axhline(y=0, color='black', linestyle='--')
    ax.set_xlabel("Lambda")
    ax.set_ylabel("Beta Estimate")
    ax.set_title("Lasso Regression Trace", fontsize=16)
    ax.legend(labels=['Room', 'Residential Zone', 'Highway Access', 'Crime Rate',
    'Tax'])
    ax.grid(True)
    plt.show()

find_alpha()

lasso = linear_model.Lasso(alpha=0.05)
lasso.fit(X_train, y_train)
print(lasso.coef_)
y_pred = lasso.predict(X_test)
print("LASSO Train score: {:.f}".format(lasso.score(X_train, y_train)))
print("LASSO Test score: {:.f}".format(lasso.score(X_test, y_test)))

plt.scatter(y_test, y_pred)
plt.show()


```



```
[ -0.83796934  0.46704064  0.          0.70483741 -1.74067035  3.16378634  
 -0.06698043 -2.67152839  1.42785917 -0.98008334 -1.90529396  1.07252204  
 -3.62223499]
```

LASSO Train score: 0.749226

LASSO Test score: 0.662584



```

In [97]: from sklearn import linear_model
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

np.random.seed(101)

#generate dataset (wide dataset)
X = np.random.normal(0,1,size=(50, 30))
# will have 10 some high coefficients
coef1 = np.random.uniform(0.5,1,10)
# and 20 low coefficeints
coef2 = np.random.uniform(0.0,0.3,20)
coef = np.concatenate((coef1, coef2), axis=0)
# plt.hist(coef)
# plt.show()
y_dot = X.dot(coef)

# generate response data by adding noise
# why adding noise?
y_response = y_dot + np.random.normal(0,1,50)

lr_train = []
lr_test = []
lr_rmse = []
ridge_train = []
ridge_test = []
ridge_rmse = []
lasso_train = []
lasso_test = []
lasso_rmse = []

for i in range(0,100):
    X_train, X_test, y_train, y_test = train_test_split(X, y_response, test_si
ze=0.2)

    ##### Linear regression #####
    reg = linear_model.LinearRegression()
    reg.fit(X_train, y_train)
    y_pred = reg.predict(X_test)
    lr_train.append(reg.score(X_train,y_train))
    lr_test.append(reg.score(X_test,y_test))
    lr_rmse.append(mean_squared_error(y_test, y_pred))

    ##### RIDGE regression #####
    ridge = Ridge(alpha=5.0)
    ridge.fit(X_train, y_train)
    y_pred = ridge.predict(X_test)
    ridge_train.append(ridge.score(X_train,y_train))
    ridge_test.append(ridge.score(X_test,y_test))
    ridge_rmse.append(mean_squared_error(y_test, y_pred))

    ##### LASSO regression #####

```

```
lasso = linear_model.Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
y_pred = lasso.predict(X_test)
lasso_train.append(lasso.score(X_train,y_train))
lasso_test.append(lasso.score(X_test,y_test))
lasso_rmse.append(mean_squared_error(y_test, y_pred))

print('LR train: ',sum(lr_train)/len(lr_train))
print('LR test: ',sum(lr_test)/len(lr_test))
print('LR RMSE: ',sum(lr_rmse)/len(lr_rmse))
print("")
print('Ridge train: ',sum(ridge_train)/len(ridge_train))
print('Ridge test: ',sum(ridge_test)/len(ridge_test))
print('Ridge RMSE: ',sum(ridge_rmse)/len(ridge_rmse))
print("")
print('Lasso train: ',sum(lasso_train)/len(lasso_train))
print('Lasso test: ',sum(lasso_test)/len(lasso_test))
print('Lasso RMSE: ',sum(lasso_rmse)/len(lasso_rmse))

#examine variance and bias
print('\n Hist Plots \n LR test: \n')
plt.hist(lr_test)
plt.show()
print('Ridge test: \n')
plt.hist(ridge_test)
plt.show()
print('Lasso test: \n')
plt.hist(lasso_test)
plt.show()
```


LR train: 0.9767790895554215
LR test: 0.3770297714569159
LR RMSE: 3.1986012122242817

Ridge train: 0.9545350496479551
Ridge test: 0.6537684262994671
Ridge RMSE: 2.0192459660320226


Lasso train: 0.917867089365298
Lasso test: 0.5547537632599874
Lasso RMSE: 2.592433720110876

Hist Plots


LR test:



Ridge test:



Lasso test:



ORANGE

Use the script below

https://www.youtube.com/playlist?list=PLmNPvQr9Tf-ZSDLwOzxpY-HrE0yv-8Fy&disable_polymer=true
 (https://www.youtube.com/playlist?list=PLmNPvQr9Tf-ZSDLwOzxpY-HrE0yv-8Fy&disable_polymer=true)

```
In [ ]: import numpy as np
        from Orange.data import Domain, Table

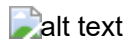
        #generate dataset (wide dataset)
        X = np.random.normal(0,1,size=(50, 30))
        # will have 10 some high coefficients
        coef1 = np.random.uniform(0.5,1,10)
        # and 20 low coefficients
        coef2 = np.random.uniform(0.0,0.3,20)
        coef = np.concatenate((coef1, coef2), axis=0)
        # plt.hist(coef)
        # plt.show()
        y_dot = X.dot(coef)
        y_response = y_dot + np.random.normal(0,1,50)

        y_response = y_response.reshape((y_response.shape[0],1))
        all_data = np.append(X, y_response, 1)
        out_data = Table(all_data)
```

Excercises

1. Add seed to both Jupyter code and to ORANGE code: `np.random.seed(101)`
2. Why are the results not the same?
3. Try generating larger data set (50, 30) -> (500, 30). What happens?

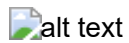
Regression Lab



Data: /data/dow_jones_index/dow_jones_index2.csv

Goal: predict closing price for the last 3 days for each stock. Date is represented as numeric value here, so you need to predict for date > 40700

Classification Lab



Data: /data/titanic.csv

Goal: Randomly split into train and test and predict "survived" variable