

Python and Data Science

Chris Jermaine

Rice University

Python

- Old language, first appeared in 1991
 - ▷ But updated often over the years
- Important characteristics
 - ▷ Interpreted
 - ▷ Dynamically-typed
 - ▷ High level
 - ▷ Multi-paradigm (imperative, functional, OO)
 - ▷ Generally compact, readable, easy-to-use
- Boom in popularity last five years
 - ▷ Now the first PL learned in many CS departments

Python: Why So Popular for Data Science?

- Dynamic typing/interpreted
 - ▷ Type a command, get a result
 - ▷ No need for compile/execute/debug cycle
- Quite high-level: easy for non-CS people to pick up
 - ▷ Statisticians, mathematicians, physicists...
- More of a general-purpose PL than R
 - ▷ More reasonable target for larger applications
 - ▷ More reasonable as API for platforms such as Spark
- Can be used as lightweight wrapper on efficient numerical codes
 - ▷ Unlike Java, for example

Python Basics

- Since Python is interpreted, can just fire up Python shell
 - ▷ Then start typing

- A first Python program

```
def Factorial (n):  
    if n == 1 or n == 0:  
        return 1  
    else:  
        return n * Factorial (n - 1)
```

```
Factorial (12)
```

Python Basics Continued

- Spacing and indentation
 - ▷ Indentation important
 - ▷ No begin/end nor
 - ▷ Indentation signals code block
- Variables
 - ▷ No declaration
 - ▷ All type checking dynamic
 - ▷ Just use

Python Basics Continued

- Dictionaries

- ▷ Standard container type is dictionary/map
- ▷ Example: `wordsInDoc = {}` creates empty dictionary

- Adding Data

- ▷ Add data by saying `wordsInDoc[23] = 16`
- ▷ Now can write something like `if wordsInDoc[23] == 16: ...`
- ▷ What if `wordsInDoc[23]` is not there? Will crash
- ▷ Protect with `if wordsInDoc.get(23, 0) .. a.` returns 0 if key 23 not defined

Encapsulation

- Functions/Procedures
 - ▷ Defined using `def myFunc (arg1, arg2) :`
 - ▷ Make sure to indent!
 - ▷ Procedure: no return statement
 - ▷ Function: return statement
- Remember:
 - ▷ No marker to end func/proc
 - ▷ It ends when you stop indenting

Loops

- Several common forms
- Looping through a range of values
 - ▷ Of form `for var in range (0, 50)`
 - ▷ Loops `for var in {0, 1, ..., 49}`
- Looping through data structures
 - ▷ Example: `for var in dataStruct`
 - ▷ loops through each entry in `dataStruct`
 - ▷ `dataStruct` can be an array, or a dictionary
 - ▷ If array, you loop through the entries
 - ▷ If dictionary, you loop through the keys

Loops Continued

- An example

```
a = {}  
a[1] = 'this'  
a[2] = 'that'  
a[3] = 'other'  
for b in a:  
    a[b]
```

```
'this'  
'that'  
'other'
```

NumPy

- NumPy is a Python package
- Most important one for data science!
 - ▷ Can use it to do super-fast math, statistics
 - ▷ Most basic type is NumPy array
 - ▷ Used to store vectors, matrices, tensors
- You will get some reasonable experience with NumPy
- Load with `import numpy as np`

NumPy Arrays: What Are They?

- Multi-dimensional array data structure
- And associated API
- Widely used for data intensive programming...
 - ▷ Linear algebra
 - ▷ Data science
 - ▷ ML

NumPy Arrays: Your Best Friend In DS

- Writing control flow code in DS programming is BAD

- ▷ Kind of like in SQL

- Python is interpreted

- ▷ Time for each statement execution generally large

- ▷ And in DS, you have a lot of data

- ▷ So this code can take a long time:

```
for b in range(0, BIG):  
    a[b] = b
```

```
sum = 0  
for b in a:  
    sum += a[b]
```

- Fewer statements executed, even if work same...

- ▷ ...means better performance!

To Reduce Number of Statements...

- Use NumPy arrays where possible
- Goal: one line of Python to process entire array!
- Some guidelines:
 - ▷ Try to replace dictionaries with NumPy arrays
 - ▷ Try to replace loops with bulk array operations
 - ▷ Backed by efficient, low-level implementations
 - ▷ This is known as “vectorized” programming

Creating and Filling NumPy Arrays

- To create a 2 by 5 array, filled with 3.14

```
>>> np.full((2, 5), 3.14)

array([[ 3.14,  3.14,  3.14,  3.14,  3.14],
       [ 3.14,  3.14,  3.14,  3.14,  3.14]])
```

- To create a 2 by 5 array, filled with 0

```
>>> np.zeros((2, 5))

array([[ 0.,  0.,  0.,  0.,  0.],
       [ 0.,  0.,  0.,  0.,  0.]])
```

More Complicated Creation Examples

- To create an array with odd numbers thru 10

```
>>> np.arange(1, 11, 2)
```

```
array([1, 3, 5, 7, 9])
```

- To “tile” an array

```
>>> np.tile (np.arange(1, 11, 2), (1, 2))
```

```
array([[1, 3, 5, 7, 9, 1, 3, 5, 7, 9]])
```

```
>>> np.tile (np.arange(1, 11, 2), (2, 1))
```

```
array([[1, 3, 5, 7, 9],  
       [1, 3, 5, 7, 9]])
```

Accessing Subparts of Arrays

- First we create a 2-d array (matrix)

```
>>> a1 = np.arange(1, 6, 1)
>>> a2 = np.arange(2, 7, 1)
>>> a3 = np.arange(3, 8, 1)
>>> a = np.row_stack ((a1, a2, a3))
>>> a
```

```
array([[1, 2, 3, 4, 5],
       [2, 3, 4, 5, 6],
       [3, 4, 5, 6, 7]])
```


Accessing Subparts of Arrays (cont)

correction: "last two rows"

- Say we want first two rows:

```
>>> a[1:, ]
```

```
array([[2, 3, 4, 5, 6],  
       [3, 4, 5, 6, 7]])
```

```
>>> a[1:]
```

```
array([[2, 3, 4, 5, 6],  
       [3, 4, 5, 6, 7]])
```

- Why does this work?
- Gets rows 1, 2, 3, and so on

Accessing Subparts of Arrays (cont)

- Say we want the last row:

```
>>> a[2:3,]  
array([[3, 4, 5, 6, 7]])
```

```
>>> a[2:3]  
array([[3, 4, 5, 6, 7]])
```

- Note: still a 2-d array. Want a vector?

```
>>> a[2:3][0]  
  
array([3, 4, 5, 6, 7])
```

Accessing Subparts of Arrays (cont)

- Now we want the second, third columns:

```
>>> a[:,1:3]
```

```
array([[2, 3],  
       [3, 4],  
       [4, 5]])
```

```
>>> a[:,np.array((1,2))]
```

```
array([[2, 3],  
       [3, 4],  
       [4, 5]])
```

- Works because `np.array((1,2))` is the array `[1, 2]`
- `a[:,np.array((1,2))]` gives you all rows, columns 1, 2

Aggregations Over Arrays

- In statistical/data analytics programming...
 - ▷ Tabulations: max, min, etc. over NumPy arrays are ubiquitous

- Key operation allowing this is sum

```
>>> a = np.arange(1, 6, 1)
```

```
>>> a
```

```
array([1, 2, 3, 4, 5])
```

```
>>> a.sum ()
```

```
15
```

Aggregations Over Arrays (cont)

- Can sum along dimension(s) of higher-d array

```
>>> a
```

```
array([[1, 2, 3, 4, 5],  
       [2, 3, 4, 5, 6],  
       [3, 4, 5, 6, 7]])
```

```
>>> a.sum (0)
```

```
array([6, 9, 12, 15, 18])
```

```
>>> a.sum (1)
```

```
array([15, 20, 25])
```

Aggregations Over Arrays (cont)

- Can find the maximum the same way

```
>>> a  
array([[10,  2,  3,  4,  5],  
       [ 2,  3, 13,  5,  6],  
       [ 3,  4,  5,  6,  7]])
```

```
>>> a.max ()
```

```
13
```

```
>>> a.max (0)
```

```
array([10,  4, 13,  6,  7])
```

```
>>> a.max (1)
```

```
array([10, 13,  7])
```

Aggregations Over Arrays (cont)

- Can find the position of the max as well

```
>>> a
```

```
array([[10,  2,  3,  4,  5],  
       [ 2,  3, 13,  5,  6],  
       [ 3,  4,  5,  6,  7]])
```

```
>>> a.argmax ()
```

```
7
```

```
>>> a.argmax (1)
```

```
array([0, 2, 4])
```

Now We Need a “Real Life” Set of Problems

- ...where we can apply some of these ideas

Latent Dirichlet Allocation (LDA)

- We will use data created by a statistical model called “LDA”
- LDA: stochastic model for generating a document corpus
- Most widely-used “topic model”
- A “topic” is a set of words that appear together with high prob
 - ▷ Intuitively: set of words that all have to do with the same subject

LDA Typically Used To Analyze Text

- Idea

- ▷ If you can analyze a corpus...
- ▷ And figure out a set of k topics...
- ▷ As well as how prevalent each topic is in each document
- ▷ You then know a lot about the corpus
- ▷ Ex: can use this prevalence info to search the corpus
- ▷ Two docs have similar topic compositions? Then they are similar!

Forward vs. Backward modeling

- Often, we want to “learn” an LDA model from an existing corpus
 - ▷ That is, you have a real data set
 - ▷ And you analyze the the data set
 - ▷ Goal: figure out how LDA model could have produced it
 - ▷ This is “backward” modeling
- But can also use it to generate a corpus
 - ▷ “Forward” modeling using LDA far less common
 - ▷ But we’ll use the forward LDA process to generate our lab data

Dictionary Models

- LDA is a “Bag of Words” model
- Does not impose ordering on words in doc
- Uses a dictionary
 - ▷ Dictionary is a map from each of m unique words in corpus
 - ▷ To a number from $\{1 \dots m\}$
- Example:
 - ▷ Dictionary might be: (0, bad) (1, I) (2, can't) (3, stand) (4, COMP101), (5, to) (6, leave) (7, love) (8, beer) (9, humanities) (10, classes)

From Dictionary to Bag of Words

- Document is a vector x

▷ $x[i]$ (i th entry in vector) is number of times dictionary word i appears in doc

0 1 2 3 4 5 6 7 8 9 10...

- Recall our dictionary is (0, bad) (1, I) (2, can't) (3, stand) (4, COMP101), (5, to) (6, leave) (7, love) (8, beer) (9, humanities) (10, classes)

- Then

freq for dict

0 1 2 3 4 5 6 7 8 9 10

▷ Sentence “I can’t stand bad beer” is $\langle 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0 \rangle$

▷ Sentence “I love beer I can’t love humanities classes” is $\langle 0, 2, 1, 0, 0, 0, 0, 2, 1, 1, 1 \rangle$

0 1 2 3 4 5 6 7 8 9 10

without order you have smaller dictionary but lose semantics, but you get back using k-grams.

LDA Step One

This example is used for generating corpus, assuming you already have a model. "Forward"

- Generate a list of the k “topics”
 - ▷ Each topic is represented by a vector of probabilities
 - ▷ $wordsInTopic_t[w]$ is the probability that topic t would produce word w
 - ▷ $wordsInTopic_t$ is sampled from a Dirichlet (α) distribution
- Example, $k = 3$
 - ▷ $wordsInTopic_0 = \langle .2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0 \rangle$
 - ▷ $wordsInTopic_1 = \langle 0, .2, .2, .2, 0, 0, 0, 0, 0, .2, .2 \rangle$
 - ▷ $wordsInTopic_2 = \langle 0, .2, .2, 0, .2, 0, .2, .2, 0, 0, 0 \rangle$

vector of probs so adds up to 1

LDA Step Two

- Generate the topic proportions for each document
 - ▷ Each topic “controls” production of some of the words in a doc
 - ▷ $topicsInDoc_d[t]$ is the probability that an arbitrary word in document d will be controlled by topic t
 - ▷ $topicsInDoc_d$ is sampled from a Dirichlet (β) distribution

LDA Step Three

- Generate the bag of words in each document
- $wordsInDoc_d[w]$ is the number of occurrences of word w in document d
- To get this vector, generate the words one-at-a-time
- For each word in d
 - ▷ Figure out the topic t that controls it:
 - ▷ Sample from a Multinomial ($topicsInDoc_d, 1$) distribution
 - ▷ Generate the word w by sampling from a Multinomial ($wordsInTopic_t, 1$) dist
 - ▷ Then increment $wordsInDoc_d[w]$

toss ball to different sized buckets topic, then toss ball into different sized buckets for word within topic

Example

- $topicsInDoc_0 = \langle .98, 0.01, 0.01 \rangle$ topic 0 is most important
- Generate first word:
 - ▷ We get $\langle 1, 0, 0 \rangle$ from a Multinomial ($topicsInDoc_0, 1$) dist
 - ▷ So we generate the word using $wordsInTopic_0 = \langle .2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0 \rangle$
 - ▷ And we get $\langle 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0 \rangle$, which is equivalent to “I”
- Generate the second word:
 - ▷ We get $\langle 1, 0, 0 \rangle$ from a Multinomial ($topicsInDoc_0, 1$) dist
 - ▷ So we generate the word using $wordsInTopic_0 = \langle .2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0 \rangle$
 - ▷ And we get $\langle 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0 \rangle$, which is equivalent to “can’t”
- Generate the third word:
 - ▷ We get $\langle 1, 0, 0 \rangle$ from a Multinomial ($topicsInDoc_0, 1$) dist
 - ▷ So we generate the word using $wordsInTopic_0 = \langle .2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0 \rangle$
 - ▷ And we get $\langle 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0 \rangle$, which is equivalent to “stand”

generating is sequential but end result is bag of words so order is not important

And the Doc Generated...

- Recall the three words generated were:
 - ▷ $\langle 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0 \rangle$
 - ▷ $\langle 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0 \rangle$
 - ▷ $\langle 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0 \rangle$
- Doc so far is $\langle 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0 \rangle$
- Keep going and get $wordsInDoc_0 = \langle 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0 \rangle$
- Encodes “I can’t stand bad beer”
~ I stand bad beer can't

OK, Back To Python!

- In a series of labs, we will look at some code that implements LDA
- Uses lots of NumPy functionality
 - ▷ `np.random.multinomial (numTrials, probVector, numRows)`
 - ▷ Take numRows samples from a Multinomial (probVector, numTrials) dist
 - ▷ Put in a matrix with numRows rows
 - ▷ `np.flatnonzero (array)`
 - ▷ Return array of indices of non-zero elements of array
 - ▷ `np.random.dirichlet (paramVector, numRows)`
 - ▷ Take numRows samples from a Dirichlet (paramVector) dist
 - ▷ `np.full (numEntries, val)`
 - ▷ Create a NumPy array with the spec'ed number of entries, all set to val

Questions?

First Activity: LDA

- Can you complete the activity?

▷ cmj4.web.rice.edu/LDADictionaryBased.html

Problem: Bad Code!

- As we said: Don't write statistical/math Python code this way
- Vectorized is better!
- Can you complete the activity?
 - ▷ `cmj4.web.rice.edu/LDAArrays.html`
 - ▷ No dictionaries here! Just arrays

Computing Cross-Tabulations

- Now that we have an array-based LDA code...
- Let's practice doing cross-tabulations on it
- Can you complete the activity?
 - ▶ <http://cmj4.web.rice.edu/Subarrays.html>

Now We'll Implement Co-Occurrence Analysis

- Fundamental task in many statistical/data mining computations
- In text processing...
 - ▷ Given a document corpus
 - ▷ Want to count number of times $(word_1, word_2)$ occur in same doc in corpus
- Your task in next activity:
 - ▷ Build three implementations
 - ▷ Utilizing varying degrees of vectorization
 - ▷ We will time each, see which is faster

Implementation One

- Nested loops
- Loop through each doc...
 - ▷ For each doc, consider each (word, word) pair it contains
 - ▷ And increment the count
- Has advantage when *wordsInCorpus* is sparse
 - ▷ Only $numDocs \times (numDistinctWordsPerDoc)^2$ execs of inner loop
- But not great in an interpreted language

Implementation Two

- Vector-based, with a loop over docs

- Given a 1-d array...


- ▷ The outer product of array with itself creates a 2-d matrix
- ▷ Where i th row is $array[i] \times array$
- ▷ So if an array gives number of occurs of each word in a doc...
- ▷ And we clip array so $\langle 0, 0, 3, 1, 0, 1... \rangle$ becomes $\langle 0, 0, 1, 1, 0, 1... \rangle$...
- ▷ Then take outer product of array with itself...
- ▷ Entry at pos (i, j) is number of co-occurs of dictionary words i, j in doc

outer product
<1, 4, 7>
*
<1, 4, 7>
1*1 1*4 1*7
1*4 4*4 7*4
etc

- Note:

- ▷ `np.outer (arrayOne, arrayTwo)` is outer product of arrays
- ▷ `np.clip (array, low, high)` clips all entries to max of high, min of low

Implementation Three



- Pure vector-based
- Note that after matrix multiply...
 - ▷ Entry at pos (i, j) is inner product of row i from LHS, col j from RHS
 - ▷ So if row i is number of occurs of word i in every doc
 - ▷ And if col j is number of occurs of word j in every doc
 - ▷ Entry at pos (i, j) is number of co-occurs of words i, j
 - ▷ Suggests a super-efficient algorithm

Implementation Three (cont)

- Some useful routines:
 - ▷ `np.transpose (array)` computes transpose of matrix in array
 - ▷ `np.dot (array1, array2)` computes dot product of 1-d arrays, matrix multiply of 2-d

These Three Implementations: The Next Activity

- Compare the three different implementations

▷ <http://cmj4.web.rice.edu/CoOccur.html>

Questions?